

## CHAPTER 11

### Training Deep Neural Networks

Training very deep neural networks introduces challenges that do not usually appear in shallow models. As networks grow deeper and wider, they become harder to optimize, more prone to overfitting, and slower to train. This chapter focuses on the practical difficulties of deep learning and presents a collection of techniques that make training deep neural networks stable, efficient, and reliable.

---

### The Vanishing and Exploding Gradients Problems

Deep neural networks are typically trained using backpropagation, which propagates error gradients from the output layer back toward the input layer. In very deep networks, these gradients often become unstable. In the **vanishing gradients problem**, gradients shrink exponentially as they flow backward, leaving the lower layers almost unchanged and preventing effective learning. In contrast, the **exploding gradients problem** occurs when gradients grow uncontrollably large, causing weight updates to diverge. Together, these issues lead to unstable training and were one of the main reasons deep neural networks were largely abandoned before 2010.

This behavior was analyzed in detail by Glorot and Bengio, who showed that the combination of poor weight initialization and saturating activation functions, such as the logistic sigmoid, amplifies variance as signals propagate through the network. As activations saturate near 0 or 1, their derivatives approach zero, leaving almost no gradient to propagate backward, especially in lower layers.




Figure 11-1. Logistic activation function saturation

---

## Glorot (Xavier) and He Initialization

To stabilize signal propagation, Glorot and Bengio proposed initializing weights so that the variance of outputs remains approximately equal to the variance of inputs, both during forward propagation and backpropagation. This idea led to **Glorot (Xavier) initialization**, which scales initial weights based on the average of a layer's fan-in and fan-out. This approach significantly reduces gradient instability and speeds up convergence.

Later work showed that different activation functions benefit from different scaling rules. In particular, **He initialization**, which scales weights based only on fan-in, works especially well for ReLU-based networks. These initialization strategies are now defaults in modern deep learning frameworks.

Initialization	Activation functions	$\sigma^2$ (Normal)
Glorot	None, tanh, logistic, softmax	$1 / \text{fan}_{\text{avg}}$
He	ReLU and variants	$2 / \text{fan}_{\text{in}}$
LeCun	SELU	$1 / \text{fan}_{\text{in}}$

Table 11-1. Initialization parameters for each type of activation function

---

## Nonsaturating Activation Functions

Activation functions play a critical role in gradient stability. Saturating activations such as sigmoid and tanh often exacerbate vanishing gradients, while nonsaturating functions improve gradient flow. The **ReLU** activation function became popular because it is simple, fast, and avoids saturation for positive inputs. However, ReLU suffers from the *dying ReLU* problem, where neurons become permanently inactive.

Variants such as **Leaky ReLU**, **PReLU**, and **RReLU** address this issue by allowing a small, nonzero gradient for negative inputs. More advanced functions like **ELU** further improve convergence by producing outputs closer to zero mean and maintaining nonzero gradients everywhere.




Figure 11-2. Leaky ReLU activation function




Figure 11-3. ELU activation function

The **SELU** activation function goes one step further by enabling *self-normalizing networks*. When combined with standardized inputs, LeCun initialization, and a purely sequential architecture, SELU preserves mean and variance across layers, effectively mitigating vanishing and exploding gradients without Batch Normalization.

## Batch Normalization

Even with good initialization and activation functions, gradients can become unstable during training. **Batch Normalization (BN)** addresses this by normalizing layer inputs during training, then learning optimal scaling and shifting parameters. By reducing internal covariate shift, BN stabilizes training, allows larger learning rates, and accelerates convergence.

BN computes batch-wise statistics during training and relies on moving averages during inference. In practice, BN often reduces or even eliminates the need for careful initialization and acts as a strong regularizer.

1.  $\mu_B = \frac{1}{m_B} \sum_{i=1}^{m_B} \mathbf{x}^{(i)}$
2.  $\sigma_B^2 = \frac{1}{m_B} \sum_{i=1}^{m_B} (\mathbf{x}^{(i)} - \mu_B)^2$
3.  $\hat{\mathbf{x}}^{(i)} = \frac{\mathbf{x}^{(i)} - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}}$
4.  $\mathbf{z}^{(i)} = \gamma \otimes \hat{\mathbf{x}}^{(i)} + \beta$

/Equation 11-3. Batch Normalization algorithm

Although BN increases training-time computation, its faster convergence often leads to shorter overall training time. BN layers can also be fused into preceding layers after training to remove runtime overhead.




Figure 11-4. Reusing pretrained layers

## Gradient Clipping

While Batch Normalization is effective for feedforward networks, **gradient clipping** is commonly used to control exploding gradients, especially in recurrent neural networks. This technique limits gradient magnitude either by value or by norm, preventing destabilizing updates while preserving useful gradient directions.

## Reusing Pretrained Layers (Transfer Learning)

Training large networks from scratch is rarely optimal. **Transfer learning** reuses layers from a pretrained model that solves a similar task, significantly reducing training time and data requirements. Lower layers typically learn general features and are more reusable than higher layers, which capture task-specific patterns.

The standard approach is to freeze reused layers initially, train new layers on top, then optionally unfreeze some upper layers and fine-tune them using a smaller learning rate. This strategy is particularly effective for deep convolutional networks.




Figure 11-5. Unsupervised pretraining workflow

## Unsupervised and Auxiliary-Task Pretraining

When labeled data is scarce and no suitable pretrained model exists, **unsupervised pretraining** can be used. Models such as autoencoders or GANs are trained on large amounts of unlabeled data to learn useful representations, which are then fine-tuned using labeled data.

A related approach is **self-supervised or auxiliary-task pretraining**, where labels are automatically generated from the data itself. This technique is especially powerful in domains such as natural language processing and computer vision, where massive unlabeled datasets are available.

## Faster Optimizers

Deep networks can take an extremely long time to train using plain Gradient Descent. Advanced optimizers dramatically improve convergence speed and stability. **Momentum optimization** accelerates training by accumulating past gradients, while **Nesterov Accelerated Gradient** further improves convergence by looking ahead in the direction of momentum.




Figure 11-6. Regular vs. Nesterov momentum

Adaptive learning-rate methods such as **AdaGrad**, **RMSProp**, **Adam**, and **Nadam** adjust learning rates dynamically for each parameter. Among these, Adam is widely used due to its robustness and minimal tuning requirements, though in some cases simpler optimizers like Nesterov momentum may generalize better.

Class	Convergence speed	Convergence quality
SGD	*	***
SGD(momentum=...)	**	***
SGD(momentum=..., nesterov=True)	**	***
Adagrad	***	* (stops too early)
RMSprop	***	** or ***
Adam	***	** or ***
Nadam	***	** or ***
AdaMax	***	** or ***

Table 11-2. Optimizer comparison

## Learning Rate Scheduling

Choosing an appropriate learning rate is crucial. Static learning rates are often suboptimal, so **learning rate schedules** are commonly used to accelerate convergence. Popular strategies include power scheduling, exponential decay, performance-based scheduling, and the **1cycle policy**, which dynamically increases and then decreases the learning rate to reach better solutions faster.




Figure 11-8. Learning curves for various learning rates

## Avoiding Overfitting Through Regularization

Deep neural networks have enormous capacity and easily overfit. In addition to early stopping and Batch Normalization, explicit regularization techniques are often necessary. **L1 and L2 regularization** constrain weight magnitudes, with L1 encouraging sparsity. **Dropout** randomly disables neurons during training, preventing co-adaptation and improving generalization.




Figure 11-9. Dropout regularization

Dropout can be extended into **Monte Carlo Dropout**, which enables uncertainty estimation at inference time by keeping dropout active and averaging multiple predictions. This approach provides more realistic confidence estimates and often improves predictive performance.

Another effective method is **max-norm regularization**, which constrains the norm of incoming weights for each neuron, further stabilizing training and reducing overfitting.

---

## Summary and Practical Guidelines

This chapter presented a comprehensive toolkit for training deep neural networks effectively. By combining proper initialization, robust activation functions, Batch Normalization, modern optimizers, learning rate schedules, and regularization techniques, it is possible to train very deep models reliably. While no single configuration fits all tasks, the practical guidelines presented provide strong defaults for most applications. With these tools, deep learning becomes not only feasible but highly effective using Keras and TensorFlow.